



Using lex/yacc to parse configuration files

One of the (many) uses of `lex` and `yacc` is to parse configuration files. A configuration file is a text file containing a set of options for a program, laid out in a very specific format. This format is usually complex enough that using `sed`, `awk`, or plain C would involve far too much code. `lex/yacc` provide an easy way of interpreting that format and, furthermore, of including the parsing code in other C programs.

In this lab, we will use an ad hoc configuration file format similar to that used by the BIND DNS server. We will start by building a simple parser that accepts valid files and rejects files that are improperly formatted (for extra sanity, this simple parser will also print some of the file's content to standard output). Then, we will progressively add more features until we end up with a powerful configuration file parser.

Exercise 1

We wish to parse configuration files such as this one:

```
global {
    num_daemons = 4
    max_out_bandwidth = 10.0
    max_in_bandwidth = 5.5
    etc = /etc/ftpd
    hostkey = /etc/ftpd/hostkey.pem
};

host "ftp-1.foobar.com" {
    ftproot = /var/ftp/server1
    max_out_bandwidth = 20.7
};

host "ftp-2.foobar.com" {
    ftproot = /var/ftp/server2
    exclude = /var/ftp/server2/private
};
```

We are given the following informal description:

- This configuration file specifies options for an FTP server. The options can be global (affecting the entire server) or host-specific.
- The file must always contain a global section, even if it is empty. This section always comes first.
- The file can contain zero or several host sections. These sections, in turn, can be empty.
- Options are specified by an *option identifier*, followed by the equals sign, followed by the *option value*. Option identifiers are alphanumerical strings (the underscore character is also allowed) beginning with a letter. Option values can be integers, real numbers, or strings. String option values can contain letters, numbers, the



underscore `_` character, the hyphen `-` character, the slash `/` character, and the dot `.` character, and must begin with a letter or a slash character.

- Host names (in the host sections) can contain the same characters as option values.
- The rest should be easy to infer from the configuration file.

You are provided with a set of files that your parser should accept (`server.conf.good.tgz`), along with a set of files that it should *not* accept (`server.conf.bad.tgz`). Do not assume anything about the configuration file format: if there are any ambiguities, ask the instructor.

If your parser accepts the file, it should print the file's content to standard output like this:

```
GLOBAL CONFIG
-----
num_daemons=4
max_out_bandwidth=10.0
max_in_bandwidth=5.5
etc=/etc/ftpd
hostcert=/etc/ftpd/hostkey.pem

HOST 'ftp-1.foobar.com' CONFIG
-----
ftproot=/var/ftp/server1
max_out_bandwidth=20.7

HOST 'ftp-2.foobar.com' CONFIG
-----
ftproot=/var/ftp/server2
exclude=/var/ftp/server2/private
```

(This output corresponds to the 'good' file `server.conf`)

If the parser rejects the file, you are allowed to print out the contents of the file up to the point when the error is detected (we will improve this in Exercise 3).

Note: This exercise is meant to be done using both `yacc` and `lex`. You should use `lex` to first tokenize the file (the informal description of the configuration file indirectly suggests what the different tokens should be: option identifier, option value...). Although you could also express the format of the file using only a `yacc` grammar, that would be overkill. Using `lex` to write the lexer (`yylex()`) makes things much easier.

Exercise 2



In this exercise, we will make the information parsed from the file available to a C program. We will do this by loading the information into a set of data structures in memory, instead of simply printing out the information to standard output in each of the parser's actions.

These data structures have already been designed for you, and are defined in `configfile.h`:

```
struct configoption {
    char *name;
    char *value;
};

typedef struct configoption configoption;

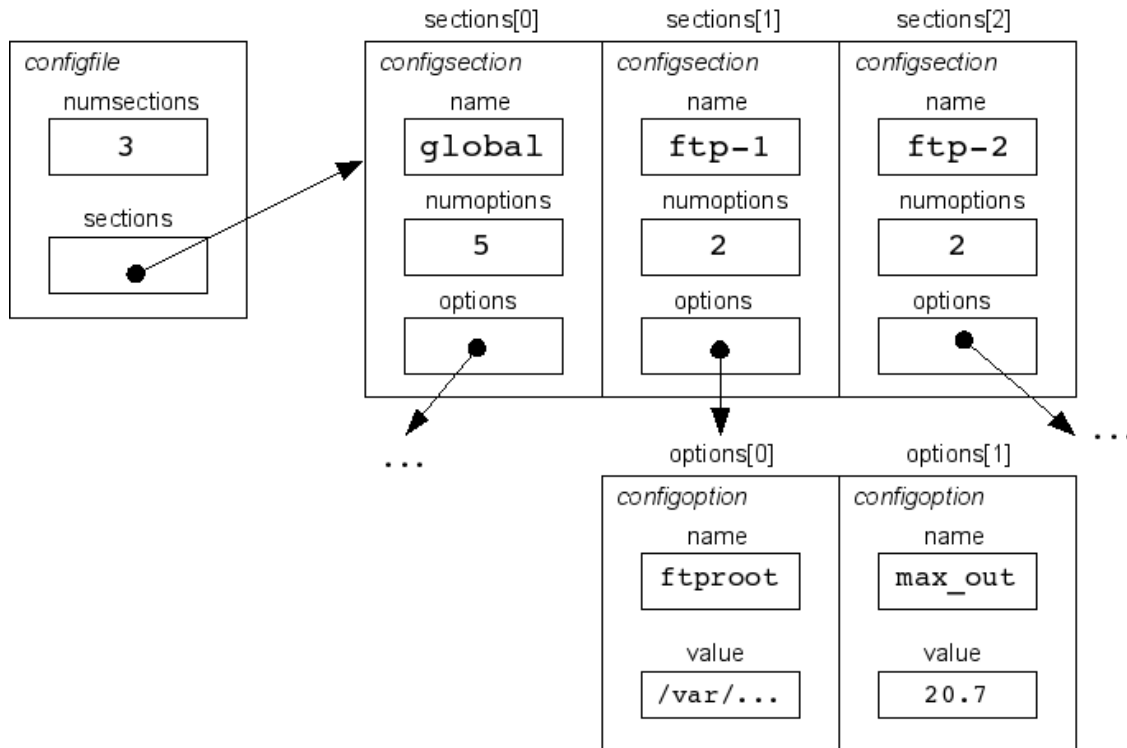
struct configsection {
    char *name;
    unsigned int numopts;
    configoption *options;
};

typedef struct configsection configsection;

struct configfile {
    unsigned int numsections;
    configsection *sections;
};

typedef struct configfile configfile;
```

You are *not* allowed to modify the specification of these structures. To give you a better idea of how you are intended to use them, the following diagram shows how a configuration file could be represented in memory:



Once the parsing is done (i.e., `yyparse()` has been called), the configuration file should be contained in a `configfile` variable. For example, your main program could look like:

```
configfile *cf;

int main (void)
{
    cf = malloc(sizeof(configfile));
    // Initialize values in cf
    yyparse ();
    // Do something with variable cf
}
```

To test your program, you will use the `cf` variable to write out the values of the options. The output of your program should be something like this:

SECTION global



```
-----  
num_daemons=4  
max_out_bandwidth=10.0  
max_in_bandwidth=5.5  
etc=/etc/ftpd  
hostcert=/etc/ftpd/hostkey.pem  
  
SECTION ftp-1.foobar.com  
-----  
ftproot=/var/ftp/server1  
max_out_bandwidth=20.7  
  
SECTION ftp-2.foobar.com  
-----  
ftproot=/var/ftp/server2  
exclude=/var/ftp/server2/private
```

Note: To make the implementation simpler, we suggest you implement the following functions, and simply make calls to them from your parsing actions:

```
// Initialize a configfile structure (0 sections)  
void init_file(configfile *cf);  
// Initialize a configsection structure with a name (and 0 options)  
void init_sec(configsection *sec, char *name);  
// Add a section to a configfile  
void add_sec(configfile *cf, configsection *sec);  
// Initialize a configoption structure with a given name and value  
void init_opt(configoption *opt, char *name, char *value);  
// Add an option to a configsection  
void add_opt(configsection *sec, configoption *opt);  
// Print the contents of a configfile structure  
void print_cf(configfile *cf);
```

Extra Credit

Both the parser and the program treat all option values as strings (type `char*`). Modify your parser and the provided data structures so integers are represented internally as `int`'s, real numbers are represented as `double`'s, and strings are represented as `char*`'s.

Exercise 3

Modify the code from Exercise 2 so your parser will be able to handle *variable substitutions*. For example, notice how the following configuration file has option values that refer to previous options:



```
global {  
    num_daemons = 4  
    max_out_bandwidth = 10.0  
    max_in_bandwidth = 5.5  
    etc = /etc/ftpd  
    hostkey = %etc%/hostkey.pem  
};  
  
host "ftp-1.foobar.com" {  
    ftproot = /var/ftp/server1  
    max_out_bandwidth = 20.7  
};  
  
host "ftp-2.foobar.com" {  
    ftproot = /var/ftp/server2  
    exclude = %ftproot%/private  
};
```

You should perform the variable substitution at the parser level, *not* at the application level. (In other words, you can't take your code from Exercise 3 and simply go through the `configfile` structure in search of substrings that begin and end with `%`.) You will need to modify your grammar to accommodate these references, and make sure that, when a new option is added to the `configfile` structure, its value has no variable references in it (i.e., you must perform all variable substitutions before storing the option value). For simplicity, we will assume that (1) you can only refer to variables within the section where the variable substitution takes place and (2) you can only refer to options that have already been parsed.



Directories, Makefiles, and executable file name

When submitting this lab, you must include two directories:

- **ex1**: This directory will contain your implementation up to Exercise 1.
- **ex23**: In Exercises 2-3 (and the extra credit question), you will build on the implementation of Exercise 1. This directory will contain your improved implementation. You should include a README file briefly describing how complete your implementation is (i.e., tell me if you implemented Exercise 2, Exercise 3, the extra credit problem, or some combination of these).

Each directory must contain a Makefile to compile the code in that directory (the Makefile must also handle all necessary calls to `lex` and `yacc`). The Makefile should have the default **Makefile** name (i.e., it should be possible to build your code by running **make** without specifying any extra parameters).

Your generated executable file must be called **configparse**.

Your submission must observe these naming conventions.